

A fully automated greedy square jigsaw puzzle solver

Dolev Pomeranz

Michal Shemesh

Ohad Ben-Shahar

dolevp@cs.bgu.ac.il

shemeshm@cs.bgu.ac.il

ben-shahar@cs.bgu.ac.il

Computer Science Department, Ben-Gurion University of The Negev, Beer Sheva, Israel

Abstract

*In the square jigsaw puzzle problem one is required to reconstruct the complete image from a set of **non-overlapping, unordered, square puzzle parts**. Here we propose a fully automatic solver for this problem, where unlike some previous work, it assumes no clues regarding parts' location and requires no prior knowledge about the original image or its simplified (e.g., lower resolution) versions. To do so, we **introduce a greedy solver which combines both informed piece placement and rearrangement of puzzle segments** to find the final solution. Among our other contributions are new compatibility metrics which better predict the chances of two given parts to be neighbors, and a novel estimation measure which evaluates the quality of puzzle solutions without the need for ground-truth information. Incorporating these contributions, our approach facilitates solutions that surpass state-of-the-art solvers on puzzles of size larger than ever attempted before.*

1. Introduction

The popular jigsaw puzzle problem exists for many centuries, long before it was first solved by a computer in the 20th century. Given n jigsaw puzzle parts of an image, it is required to reconstruct the complete image. Although phrased as a game, this problem, which was proved to be NP-complete [2, 7], serves a platform for many applications, e.g. in biological [14], molecular [22], text and speech [12, 25], archaeological contexts [3, 10], and image editing [4].

The traditional jigsaw puzzle problem assumes pieces of various shapes, and indeed earlier computational work considered the problem in this form. The first jigsaw solver proposed by Freeman and Garder in 1964 [8] was designed to solve apictorial puzzles and was able to handle nine-piece problems. By the end of the century, shaped-based solvers were already able to reconstruct puzzles of more than 200 pieces [9]. The use of appearance and chromatic information began only three decades after the work of Freeman

and Garder (e.g., [11, 6, 21, 24, 23, 13, 15]).

A major building block in most appearance-based puzzle solvers evaluate the similarity between two parts, and the strategies for doing so typically divide into two groups. One approach compares the appearance of the abutting boundaries while using a formal distance measure between these two vectors (e.g., [11, 6, 21, 24, 13, 15]). Alternatively, it was suggested to use the entire part and measure its statistical properties in order to group similar parts together [15] or as a means to define pairwise similarity measures [11]. Once the similarity between two parts is estimated, different classes of algorithms are employed to solve the puzzle, and in particular, previous methods have used **greedy algorithms [7, 15], dynamic programming [1], genetic algorithms [21] and approximation algorithms on graphical models [5]**.

Recently, Cho *et al.* [5] presented a **probabilistic solver** which achieves approximated puzzle reconstruction via a graphical model and a probability function that is maximized via loopy belief propagation. Since they lack a local evidence term required for their graphical model, they employed two strategies that exploit prior knowledge - either a dense-and-noisy evidence which estimates the low resolution image from a bag of parts, or a sparse-and-accurate evidence which assumes that a few parts, called “anchor patches”, are given by an oracle (e.g., a human observer) at their correct location in the puzzle. While only semi-automatic, their approach was seminal in its ability to handle puzzles with over 400 pieces.

In this paper we challenge the state-of-the-art in jigsaw puzzle solving in several ways. We suggest a computational framework that can handle in reasonable time square jigsaw puzzles of size larger than ever attempted before. However, **we completely exclude the use of clues, oracles, or human intervention, as well as any use of prior knowledge about the original image or simplified (e.g., lower resolution) versions of it.** Despite these significant restrictions, our approach achieves better than state-of-the-art performance, and frequently succeeds in providing completely accurate puzzle solutions.

Our puzzle solving framework is based on a greedy solver, which works in several phases. First a *compatibility function* is computed to measure the affinity between two neighboring parts (as often done in other solvers as well). Then the solver solves three problems - **the placement problem, the segmentation problem, and the shifting problem**. The placement module places all parts on the board in an informed fashion, the segmentation module identifies regions which are likely to be assembled correctly, and the shifting module relocates regions and parts to produce the final result.

Our main contribution is a fully automated solver which uses no clues, hints, or other prior knowledge whatsoever. In addition, our contributions also include the proposal of new and better compatibility metrics, as well as the introduction of the concept of an *estimation metric* which evaluates the quality of a given solution without any reference to the original, ground-truth image. As we show, these metrics become a critical tool that facilitates self evaluation in case when no clues or prior knowledge are available. In what follows we first discuss these, as well as the rest of the metrics involved in the solver.

2. Metrics

2.1. Compatibility Metrics

A compatibility metric is at the foundation of every jigsaw solver. Given two puzzle parts x_i, x_j and a possible spatial relation R between them, the compatibility function predicts the likelihood that these two parts are indeed placed as neighbors with relation R in the correct solution. With $R \in \{l, r, u, d\}$, we use $C(x_i, x_j, R)$ to denote the compatibility that part x_j is placed on either of the left, right, up, or down side of part x_i , respectively.

Optimally, one would wish to find a metric C such that $C(x_i, x_j, R) = 1$ iff part x_j is located to the R side of x_i in the original image and 0 otherwise. If such a function exists, the jigsaw puzzle problem could be solved in polynomial time by a greedy algorithm [7]. Motivated by the above, we first discuss how one can measure the accuracy of a compatibility function and then seek a compatibility function C which is as accurate and discriminative as possible.

Measuring Compatibility Accuracy: In order to compare between compatibility metrics, we denote the classification criterion as the ratio between correct placements and the total number of possible placements. Similar to Cho *et al.* [5], we define the *correct placement of parts x_i and x_j according to the compatibility metric C* if part x_j is in relation R to part x_i in the original image and if

$$\forall x_k \in \text{Parts}, \quad C(x_i, x_j, R) \geq C(x_i, x_k, R). \quad (1)$$

Recently, Cho *et al.* [5] evaluated five compatibility met-

rics, among which the **dissimilarity-based compatibility** metric showed to be the most discriminative. Inspired by both the dissimilarity-based metric and the characteristics of natural images, we propose two new types of compatibility metrics, which will be described shortly after we review the dissimilarity-based compatibility metric.

Dissimilarity-based Compatibility: The dissimilarity between two parts x_i, x_j can be measured by **summing up the squared color differences of the pixels along the parts' abutting boundaries** [5]. For example, if we represent each color image part in normalized LAB space by a $K \times K \times 3$ matrix (where K is the part width/height in pixels) then the dissimilarity between parts x_i and x_j , where x_j is to the right of x_i , can be defined as

$$D(x_i, x_j, r) = \sum_{k=1}^K \sum_{d=1}^3 (x_i(k, K, d) - x_j(k, 1, d))^2. \quad (2)$$

We do emphasize that dissimilarity is not a distance measure, i.e., $D(x_j, x_i, R)$ is not necessarily the same as (and almost always different than) $D(x_i, x_j, R)$.

Although the dissimilarity-based metric was shown to be the most discriminative among the tested metrics in Cho *et al.* [5], the observation that it is related to the L_2 norm of the boundaries' difference vector suggests that other norms of the same difference vector could behave even better. **Inspired by the use of non Euclidean norms in image operations such as noise removal [19] or diffusion [20], and by the observation that the L_2 norm penalizes large boundary differences very severely even though such large differences do exist in natural images, we evaluated other L_p norms as well.** The average accuracy empirical results for the dissimilarity metric based on various L_p norms are shown in Fig. 1 and indicate clearly that the L_2 is suboptimal, and that the best results may be obtained with L_p norms with $p \approx 0.3$. While this is outside the scope of this paper, an interesting question that emerges from this data is why does performance break down for small p values, and why does $p \approx 0.3$ appear to provide optimal performance (see the Discussion and Future Work section).




Figure 1. Comparing dissimilarity-based metrics using different L_p norms. For this test we used a database of 20 test images [5], and analyzed each for the portion of correct part pairs that received the highest compatibility score among other candidates. The plot depicts the average classification accuracy of 20 images and shows how performance peaks at $p \approx 0.3$, with a value of 86%.

$(L_p)^q$ Compatibility: As mentioned before, the optimal

compatibility function should also be as discriminative as possible. In order to reflect the scattering of the dissimilarity values when obtaining their compatibility measures, we also experimented with different powers q of the L_p norms. For example, the $(L_p)^q$ compatibility for parts x_i, x_j where x_j is to the right of x_i is defined as:

$$D_{p,q}(x_i, x_j, r) = \left(\sum_{k=1}^K \sum_{d=1}^3 (|x_i(k, K, d) - x_j(k, 1, d)|)^p \right)^{\frac{q}{p}}. \quad (3)$$

Having these dissimilarity values, we also propose to obtain its compatibility measure slightly differently than done previously. In particular, we define

$$C(x_i, x_j, R) \propto \exp \left(- \frac{D_{p,q}(x_i, x_j, R)}{\text{quartile}(i, R)} \right), \quad (4)$$

where $\text{quartile}(i, R)$ is the quartile of the dissimilarity between all other parts in relation R to part x_i . The quartile normalization gives us valuable information about the scattering of the compatibility function values, and in particular emphasizes discriminative scattering.

Although the value of q does not affect the dissimilarity classification accuracy, it does have a significant effect on our solver's performance. While Cho *et al.* [5] used $p = 2$ and $q = 2$ in their dissimilarity metric (Eq. 2), we found that optimal results are achieved with power value $q = 1/16$ and hence prefer to use a $(L_{p=3/10})^{q=1/16}$, i.e.,

$$D_{p,q}(x_i, x_j, r) = \left(\sum_{k=1}^K \sum_{d=1}^3 (|x_i(k, K, d) - x_j(k, 1, d)|)^{\frac{3}{10}} \right)^{\frac{5}{24}}. \quad (5)$$

Prediction-based Compatibility: While dissimilarity-based metrics may be improved greatly by employing $(L_p)^q$ affinities, other possibilities may be promising as well. In particular, as opposed to measure differences, one may attempt to quantify how well one can predict the boundary content of one part based on the boundary content of the other. The better the prediction, the higher the compatibility of the two parts, a measure we call a *Prediction-based Compatibility*.

Naturally, **predictions over the boundaries can be made using Taylor's expansion**. A first order prediction would need an estimation of the derivative of each part at its boundary, a computation that can be done numerically in one of several ways. For example, employing backward differences estimation using the last two pixels in each row near the boundary, one can obtain a prediction of the first pixel in the neighboring part. The quality of this prediction can then be verified against the actual pixel value from the second part by using any desired norm, and in general, using $(L_p)^q$ for preselected p and q as discussed above in the dissimilarity measure. Importantly, we repeat this computation for both parts in a symmetric fashion. For example, the prediction metric for two parts x_i, x_j with relation r when

using $(L_p)^q$ with $p = 3/10$ and $q = 1/16$ can be expressed by the following expression

$$\text{Pred}(x_i, x_j, r) = \sum_{k=1}^K \sum_{d=1}^3 \left[([2x_i(k, K, d) - x_i(k, K-1, d)] - x_j(k, 1, d))^{\frac{3}{10}} + ([2x_j(k, 1, d) - x_j(k, 2, d)] - x_i(k, K, d))^{\frac{3}{10}} \right]^{\frac{5}{24}} \quad (6)$$

and obtain its compatibility measure as in Eq. 4.

Comparison of Compatibility Metrics: Fig. 2 shows the comparison between compatibility metrics' accuracy on three different databases. Note that even though dissimilarity-based and $(L_p)^q$ metrics are directly related in definition, $(L_{3/10})^{1/16}$ has achieved a significant improvement of 7% over the dissimilarity metric (86% vs. 79%). As shown, the prediction-based compatibility performs very similarly, but further evaluation on additional puzzle databases (see Sec 4) clearly indicates its superiority.




Figure 2. Comparing compatibility metrics' accuracy. For each image in the test sets consisting of 20 images each, we find the portion of correct part pairs that received the highest compatibility score among other candidates. We show the average classification accuracy of 20 images for each test set. Note how the $(L_{3/10})^{1/16}$ and the prediction-based compatibilities perform better than the dissimilarity-based compatibility in all three sets.

Before we proceed to describe our jigsaw solver, we discuss two additional types of metrics. The first, the *performance metric*, evaluates the performance of an approximated solution by comparison to the ground truth [5]. The second, a novel metric which we term the *estimation metric*, evaluates the quality of an approximated solution *without* prior knowledge about the original image. The motivation for this kind of metric is driven from the need to assess the quality of an approximated solution in the process of computing it, when no access to ground-truth solution is available.

2.2. Performance Metrics

Performance metrics are used to evaluate the performance accuracy of an approximated solution by comparison to the correct (i.e., desired) solution. Cho *et al.* [5] introduced three performance metrics, two of which are relevant to our algorithm and are reviewed below.

Direct Comparison Metric: The direct comparison metric calculates the ratio between the number of parts in the approximated solution which are placed in their correct location and the total number of parts [5].

Neighbor Comparison Metric: The neighbor comparison metric calculates the ratio between the number of correct neighbors placement for each part and the total number of neighbors [5].

By construction, the direct comparison metric does not tolerate parts which are not placed in their correct location, and hence grants lower scores even to solutions that are assembled correctly but are slightly shifted. Performance metrics' accuracy results of Cho *et al.* and our solver are shown in Fig. 3.




Figure 3. Performance metrics' accuracy comparison. We compare our results with those reported by Cho *et al.* [5] in both the direct measure (upper graph) and neighbor measure (lower graph). Under the direct comparison measure Cho *et al.* average reconstruction accuracy is lower than 10% while ours is 94%. Under the neighbor comparison measure Cho *et al.* average reconstruction accuracy is 55% while ours is 95%.

2.3. Estimation Metrics

While the performance metrics evaluate the quality of solutions in the presence of a ground truth image, the *estimation metric* is designed to provide such quality scores without knowing the correct solution whatsoever. Given the fact that we use no clues or prior information about the original image, this is an essential tool which can be used during the reconstruction process in order to improve the solver's final approximated solution. In this paper we introduce this notion and suggest one estimation metric that is put to use in our proposed puzzle solver.

The “Best Buddies” metric: Two parts x_i, x_j , their relation R_1 and opposite relation R_2 are said to be *best buddies* iff the following holds:

$$\begin{aligned} \forall x_k \in Parts, \quad C(x_i, x_j, R_1) &\geq C(x_i, x_k, R_1) \\ \text{and} \\ \forall x_p \in Parts, \quad C(x_j, x_i, R_2) &\geq C(x_j, x_p, R_2). \end{aligned} \quad (7)$$

Intuitively, two parts are best buddies if both “agree” that the other part is their most likely neighbor in a certain spa-

tial relation. The best buddies *metric* for a given approximated solution represents the ratio between the number of neighbors who are said to be best buddies and the total number of neighbors.

Testing the correlation between the best buddies estimation metric and the performance metrics reveals no correlation to the direct performance metric (which is not based on pairwise similarities) but a strong correlation to the neighbor performance metric. Fig. 4 depicts this relationship as a scatter plot of the best buddies estimation metric vs. the neighbor performance metric. Every dot in the graph represent a single puzzle solver solution, where the x-coordinates and the y-coordinate denote the performance accuracy and estimation accuracy, respectively. **Each image in the set was tested 10 times, totaling 200 sample points.**




Figure 4. Correlation between the best buddies estimation metric and the neighbor comparison performance metric. Every dot in this 200 point graph represents a single puzzle solver solution, where the x and y coordinates denote its performance value and estimation value, respectively.

3. A greedy puzzle solver

Equipped with the different types of metrics discussed above, we are now ready to describe our puzzle solver. At its higher level, this solver divides the puzzle reconstruction problem into three subproblems, and solves each of them separately.

- The *placement problem*: Given a single part or a partial constructed puzzle, find the position of the remaining parts on the board.
- The *segmentation problem*: Given a placement of all parts on the board, i.e. an approximated solution, divide it to segments which are estimated to be assembled correctly, disregarding their absolute location.
- The *shifting problem*: Given a set of puzzle segments, relocate both segments and individual parts on the board such that a better approximated solution is obtained.

In what follows we describe algorithms for all these problems.

3.1. The Placer

As the first solver module, the *placer* is a greedy algorithm that is given a single part or a partially constructed puzzle, and places the remaining parts on the board in an informed fashion. More specifically, given a greedy criterion and a “seed” part (or a partial constructed seed puzzle), the placer reconstructs the remaining puzzle parts around the seed by iteratively applying the greedy criterion. Such greedy criteria are discussed below.

Technically, the greedy procedure is affected by the board’s dimensions such that piece placement is not allowed to exceed the puzzle’s width or height. Furthermore, our greedy criteria are also endowed with the heuristics that *empty* puzzle slots with a larger number of *occupied* slots (i.e., slots in which parts had been allocated), are more informative, and hence should be allocated a puzzle part sooner than others. Note that using this heuristic while starting from a single seed part limits empty slots to have at most two occupied neighbor slots, and hence that the evolving constructed puzzle will preserve a roughly rectangular shape. If the seed is a partially constructed puzzle, the placer would first complete it to a rectangular shape by placing new parts in empty slots with more than two occupied neighbor slots.

What greedy criteria can one use for placing parts in slots? Given the potential of the best buddies estimation metric (see Sec. 2.3), we propose to use it as a greedy placement criterion as well. By using it, **in each step the placer chooses the empty slot and unallocated part which “agree” to be best buddies among all empty slots and remaining unallocated parts.** If there is more than one possibility, or none at all, the algorithm then chooses the part with the highest compatibility metric to its candidate neighbors¹.

The results of using the best buddies metric as a greedy criterion for the placer is shown in Figs. 5a and 7a and clearly demonstrate that placements tend to contain regions which are assembled correctly, but are often misplaced relative to each other. This is expected, since a compatibility function can only predict the likelihood of immediate neighbors, rather than global arrangements or true locations on the puzzle. This is also the motivation for the two subsequent phases of our solver, as described next.

Before we move to do so, one more comment is in place. As opposed to previous work, our placement procedure is guaranteed to place each puzzle part exactly once, rather than more than once or even not at all, as could happen in Cho *et al.* [5]. This is due to the fact that in their graphical model the exclusion term in the maximized probability discourages parts from being used more than once, but does not prevent that entirely, therefore allowing the repetition or

¹In this case, if a candidate part is considered to an empty slot with more than one allocated neighbor, its compatibility is computed by averaging its individual compatibilities to each of these neighbors.

the complete dropping of some of the parts.

3.2. The Segmenter

Given a complete placement of all parts by the placer, we are now looking for regions which are assembled correctly, which we denote as *segments*. Clearly, although this would be easier if the ground truth solution image was available, here we care to consider again the stricter version of the problem where no such solution is available for inspection. Under such conditions we attempt to find these segments using a region growing segmentation algorithm [17] with random seeds and a segmentation predicate based on the best buddies metric (cf. Sec. 2.2). This means that two neighbor parts will be in the same segment only if they are best buddies. We have tested the segments defined by the segmenter and found that 99.7% of the neighbors in each segment are genuine neighbors in the original image. The results of this second solver module, the *segmenter*, are shown in Figs. 5b and 7b.

3.3. The Shifter

Once coherent solution segments are computed, we finally relocate the parts on the board such that a better puzzle solution is reconstructed. Towards that goal, the *shifter* iteratively repeats the following steps:

1. **It uses the current largest segment (S_{max}) in the most recent reconstructed puzzle as a new seed while executing the greedy placer once again.** Note that since the puzzle grows around the seed with no reference to a particular position, this is equivalent to shifting S_{max} to some better place as determined by the greedy placement procedure.
2. Once the placer completes a new puzzle reconstruction, it performs the segmentation again.

This loop is repeated until the evaluation of the best buddies metric reaches a local maximum.

4. Experimental Results

The following benchmark shows our solver performance results. To obtain truly comparative results to the state-of-the-art, we used the same database of 20 images from Cho *et al.* [5], where each puzzle problem consists of **432 parts of 28×28 pixels.** For these experiments we used a solver with the prediction-based compatibility metric. **For each puzzle problem the solver is executed 10 times, each with a random seed, and the solution selected is the one having the highest best buddies estimation metric score.** The average performance metrics accuracy results for all 20 images are 94% accuracy under the **Direct comparison**, and 95% under the **Neighbor comparison**, which is a significant improvement over past methods (lower than 10% and 55%, respectively,




Figure 5. Examples of results of our solver with the same images used by Cho *et al.* [5]. For each image i (as numbered in the original collection), panel $i - a$ shows the initial placement results of the greedy placer. Panel $i - b$ shows the segments of panels $i - a$ (depicted by different colors). Panel $i - c$ shows the final result of the shifter. In both cases the results are similar to the original image (shown on the right). Image 2 direct and neighbor reconstruction accuracy are 82% and 81%, respectively. Image 3 was constructed with 100% accuracy.

in Cho *et al.* [5]). Importantly, we were able to obtain 100% accurate solutions for 65% of the images, as can be seen in the supplementary material. The contribution of the shifting module was significant, since the direct and neighbor accuracy scores were 35% and 74% after the placer’s initial reconstruction and prior to the shifting phase. On average, it took the solver roughly 1.2 minutes (on a desktop with a 3.20 GHz CPU with 4 GB RAM) and 6 iterations to generate a solution. Fig. 5 shows our solver’s results for two images. In addition to the final result, this figure also shows the results of the initial placement by the placer, the initial segmentation by the segmenter, and the final result after the shifter.

As shown in Fig. 6, the final results of our solver are superior to the state-of-the-art (Cho *et al.* [5]) even though it *never uses externally provided clues*. Note in particular how some puzzle problems are not solved perfectly by previous methods even after being provided with 10 anchors, while our approach does so successfully with no externally provided prior whatsoever. Clearly, comparing these two methods on equal grounds (shown in columns a,d of Fig. 6) shows significant qualitative improvement.

Fig. 7 shows our solver’s results on puzzle problems of various part numbers, including sizes never attempted before. We used two additional databases of 20 images, each divided to parts of size 28×28 pixels each. The first, divided to 540 parts, was taken from McGill image database [16]. The second, divided to 805 parts, was self created (all images can be inspected in the supplementary material). The average reconstruction accuracy for the direct and neighbor comparison is 83% and 91% respectively for the 540 parts database, while 50% of the images are constructed

with 100% accuracy. Here it took an average of 1.9 minutes and 6.5 iterations to generate a solution. The average reconstruction accuracy for the direct and neighbor comparison is 80% and 90% respectively for the 805 parts database, while 35% of the images are constructed with 100% accuracy. In this case solutions were generated on average after 5.1 minutes and 8 iterations. While performance is expectedly decreasing, note how it surpasses previously reported performance (e.g., [5]) for problem size as large as 432 parts. In addition, Fig. 7 shows two examples of puzzles whose size is an order of magnitude larger. An exhaustive set of results is shown in the supplementary material and we note that the jigsaw solver code and images are available for the use of the community [18].

5. Discussion and Future Work

We presented a new greedy solver for the (square) jigsaw puzzle problem by employing part placement, region segmentation, and shifting phases, all informed by a novel part compatibility and solution estimation measures. Unlike previous work, our solver uses no clues or priors about the original image or the problem solution, and it requires no manual interaction, yet it still exhibits much improved results compared to the state-of-the-art. Arguably, the ability of our solver to provide successful solutions for puzzles whose size is an order of magnitude larger than ever attempted before despite using no clues, hints, or other priors, implies that this problem may be simpler than previously thought.

Several main issues still deserve further investigation in the context of our solver and jigsaw puzzle solving in general. For example, formulating compatibility metrics that




Figure 6. Comparison of our results to the state-of-the-art. Shown are the results of Cho *et al.* [5] for both their dense-and-noisy local evidence, and sparse-and-accurate local evidence with various numbers of anchors. Note that in some cases, even 10 anchors provide inferior results to those obtained by our solver (with no externally provided prior whatsoever). Clearly, comparing these two methods on equal grounds (i.e., with no externally provided clues, shown in columns a,d), demonstrates significant qualitative improvement.

can generalize compatibilities to failure cases (e.g., around roughly homogeneous image regions; see supplementary material) would clearly achieve improved results. Furthermore, a critical issue for consideration is **the choice of initial seed**, which has great effect over the final solution. Identifying a criteria that successfully **predict a good seed** (as opposed to repeated random selections) could help reduce the solver’s computational cost. Finally, while the degradation effect of **increasing the number of parts** has been explored empirically in the puzzle solving community, much less insight has been obtained about the effect of part *size*, or how uncertainty regarding parts’ appearance could affect the solution. We believe that studying such issues will improve solvers’ accuracy significantly, and will contribute greatly to the problem of jigsaw puzzles in general.

Acknowledgments

This work was funded in part by the US Air Force European Office of Aerospace Research and Development grant number FA8655-09-1-3016, the Israel Science Foundation (ISF) grant No. 1245/08 and the European Commission in the 7th Framework Programme (CROPS GA no 246252). We also thank the generous support of the Frankel fund, the Paul Ivanier center for Robotics Research and the Zlotowski Center for Neuroscience at Ben-Gurion University.

References

- [1] N. Alajlan. Solving square jigsaw puzzles using dynamic programming and the hungarian procedure. *American Journal of Applied Sciences*, 11:1942–1948, 6 2009. 9
- [2] T. Altman. Solving the jigsaw puzzle problem in linear time. *Applied Artificial Intelligence*, 3(4):453–462, 1990. 9
- [3] B. J. Brown, C. Toler-Franklin, D. Nehab, M. Burns, D. Dobkin, A. Vlachopoulos, C. Doumas, S. Rusinkiewicz, and T. Weyrich. A system for high-volume acquisition and matching of fresco fragments: Reassembling Theran wall paintings. *ACM Transactions on Graphics*, 27(3), 2008. 9
- [4] T. S. Cho, S. Avidan, and W. T. Freeman. The patch transform. *IEEE Trans. Pattern Anal. Mach. Intell.*, 32(8):1489–1501, 2010. 9
- [5] T. S. Cho, S. Avidan, and W. T. Freeman. A probabilistic image jigsaw puzzle solver. In *Proc. CVPR*, pages 183–190, 2010. 9, 10, 11, 12, 13, 14, 15
- [6] M. G. Chung, M. M. Fleck, and D. A. Forsyth. Jigsaw puzzle solver using shape and color. In *In Proceedings of the International Conference on Signal Processing*, 1998. 9
- [7] E. Demaine and M. Demaine. Jigsaw puzzles, edge matching, and polyomino packing: Connections and complexity. *Graphs and Combinatorics*, 23:195–208, 2007. 9, 10
- [8] H. Freeman and L. Garder. Apictorial jigsaw puzzles: the computer solution of a problem in pattern recognition. *Elec-*




Figure 7. Selected results of our solver for larger puzzle problems. The first row shows a result for a puzzle containing 805 parts, solved with 100% accuracy. Please refer to the text for average solution performance on the entire database of 540 and 805 parts puzzles. The second and third examples show a single solution for puzzle of 2360 and 3300 parts, respectively, (each 28×28 pixels in size), where performance was 0.21% and 95% accurate in terms of the direct and neighbor metrics for the second example and 100% for the third.

- tronic Computers, IEEE Transactions on*, 13:118–127, 1964. 9
- [9] D. Goldberg, C. Malon, and M. W. Bern. A global approach to automatic solution of jigsaw puzzles. In *Symposium on Computational Geometry*, pages 82–87, 2002. 9
- [10] D. Koller and M. Levoy. Computer-aided reconstruction and new matches in the forma urbis romae. *Bullettino Della Commissione Archeologica Comunale di Roma*, 15:103–125, 2006. 9
- [11] D. Kosiba, P. Devaux, S. Balasubramanian, T. Gandhi, and K. Kasturi. An automatic jigsaw puzzle solver. In *Pattern Recognition*, volume 1, pages 616–618, 1994. 9
- [12] M. Levison. The computer in literary studies. In A. D. Booth, editor, *Machine Translation*, pages 173–194, 1967. 9
- [13] M. Makridis and N. Papamarkos. A new technique for solving a jigsaw puzzle. In *ICIP*, pages 2001–2004, 2006. 9
- [14] W. Marande and G. Burger. Mitochondrial dna as a genomic jigsaw puzzle. In *Science*, volume 318, page 415, 2007. 9
- [15] T. R. Nielsen, P. Drewsen, and K. Hansen. Solving jigsaw puzzles using image features. *Pattern Recognition Letters*, 29(14):1924–1933, 2008. 9
- [16] A. Olmos and F. A. A. Kingdom. McGill calibrated colour image database. <http://tabby.vision.mcgill.ca>, 2005. 14
- [17] I. Pitas. *Digital Image Processing Algorithms and Applications*. John Wiley and sons, INC., 2000. 13
- [18] D. Pomeranz, M. Shemesh, and O. Ben-Shahar. A fully automated greedy square jigsaw puzzle solver MATLAB code and images. <https://sites.google.com/site/greedyjigsawsolver/home>, 2011. 14
- [19] L. Rudin, S. Osher, and E. Fatemi. Nonlinear total variation based boise removal algorithms. *Physica D*, 60:259–268, 1992. 10
- [20] B. Tang, G. Sapiro, and V. Caselles. Diffusion of general data on non-flat manifolds via harmonic maps theory: The direction diffusion case. *Int. J. Comput. Vision*, 36(2):149–161, 2000. 10
- [21] F. Toyama, Y. Fujiki, K. Shoji, and J. Miyamichi. Assembly of puzzles using a genetic algorithm. In *ICPR*, pages 389–392, 2002. 9
- [22] C.-S. E. Wang. *Determining molecular conformation from distance or density data*. PhD thesis, MIT, 2000. 9
- [23] M. Weiss-Cohen and Y. Halevi. Knowledge retrieval for automatic solving of jigsaw puzzles. In *Computational Intelligence for Modelling, Control and Automation, 2005 and International Conference on Intelligent Agents, Web Technologies and Internet Commerce, International Conference on*, volume 2, pages 379–383, 2005. 9
- [24] F.-H. Yao and G.-F. Shao. A shape and image merging technique to solve jigsaw puzzles. *Pattern Recognition Letters*, 24(12):1819–1835, 2003. 9
- [25] Y.-X. Zhao, M.-C. Su, Z.-L. Chou, and J. Lee. A puzzle solver and its application in speech descrambling. In *Proceedings of the 2007 annual Conference on International Conference on Computer Engineering and Applications*, pages 171–176, 2007. 9