

K-Nearest Neighbor Finding Using MaxNearestDist

Hanan Samet, *Fellow, IEEE*

Abstract—Similarity searching often reduces to finding the k nearest neighbors to a query object. Finding the k nearest neighbors is achieved by applying either a depth-first or a best-first algorithm to the search hierarchy containing the data. These algorithms are generally applicable to any index based on hierarchical clustering. The idea is that the data is partitioned into clusters that are aggregated to form other clusters, with the total aggregation being represented as a tree. These algorithms have traditionally used a lower bound corresponding to the minimum distance at which a nearest neighbor can be found (termed MINDIST) to prune the search process by avoiding the processing of some of the clusters, as well as individual objects when they can be shown to be farther from the query object q than all of the current k nearest neighbors of q . An alternative pruning technique that uses an upper bound corresponding to the maximum possible distance at which a nearest neighbor is guaranteed to be found (termed MAXNEARESTDIST) is described. The MAXNEARESTDIST upper bound is adapted to enable its use for finding the k nearest neighbors instead of just the nearest neighbor (that is, $k = 1$) as in its previous uses. Both the depth-first and best-first k -nearest neighbor algorithms are modified to use MAXNEARESTDIST, which is shown to enhance both algorithms by overcoming their shortcomings. In particular, for the depth-first algorithm, the number of clusters in the search hierarchy that must be examined is not increased thereby potentially lowering its execution time, while for the best-first algorithm, the number of clusters in the search hierarchy that must be retained in the priority queue used to control the ordering of processing of the clusters is also not increased, thereby potentially lowering its storage requirements.

Index Terms— k -nearest neighbors, similarity searching, metric spaces, depth-first nearest neighbor finding, best-first nearest neighbor finding.



1 INTRODUCTION

SIMILARITY searching is an important task when trying to find patterns in applications involving mining different types of data such as images, video, time series, text documents, DNA sequences, and so forth. Similarity searching often reduces to finding the k nearest neighbors to a query object. This process is facilitated by building an index on the data, which is usually based on a hierarchical clustering. The idea is that the data objects are partitioned into clusters (termed *nonobjects*), which are aggregated to form other clusters, with the total aggregation being represented as a tree known as a search hierarchy. Numerous search hierarchies have been used for both vector data and nonvector data such as data lying in a metric space many of which are surveyed in [7], [10], [15], [20], [26], [30]. The methods that we describe in this paper are independent of the nature of the data.

The most common strategy for finding the k nearest neighbors is the depth-first method, which explores the elements of the search hierarchy in a depth-first manner (for example, [14]). The k nearest neighbors found so far are kept track of in a set L with the aid of a variable D_k that indicates the distance, using a suitably defined distance function d , of the current k th-nearest object from the query object q . The depth-first method visits every element of the search

hierarchy. The *branch and bound* variant of the depth-first method yields better performance by not visiting every nonobject and its objects when it can be determined that it is impossible for the nonobject to contain any of the k nearest neighbors of q [14], [21]). For example, this is true if we know that for every nonobject element e of the search hierarchy, $d(q, e) \leq d(q, e_0)$ for every object e_0 in e and that the relation $d(q, e) > D_k$ is satisfied.¹ This can indeed be achieved if we define $d(q, e)$ as the minimum possible distance from q to any object e_0 in nonobject e (referred to as MINDIST in contrast to MAXDIST, the maximum possible distance, which unlike MINDIST cannot be used for pruning).

Letting $A(e)$ denote the set of nonobject immediate descendants e_i of nonobject element e of the search hierarchy, using the above definition of distance for nonobject elements (that is, MINDIST) makes it possible to obtain even better performance as a result of speeding up the convergence of D_k to its final value by processing elements of $A(e)$ in increasing order of $d(q, e_i)$ (that is, a MINDIST ordering). In this way, once an element e_i in $A(e)$ is found such that $d(q, e_i) > D_k$, then $d(q, e_j) > D_k$ for all remaining elements e_j of $A(e)$. This means that none of these remaining nonobject descendants of e need to be processed further, and the algorithm backtracks to the parent of e or terminates if e is the root of the search hierarchy.

An alternative strategy is the best-first method (for example, [1], [8], [11], [16], [17], [18], [23]), which explores the nonobject elements of the search hierarchy in increasing order of their distance from q (hence the characterization as

• The author is with the Computer Science Department, Center for Automation Research, Institute for Advanced Computer Studies, University of Maryland, 4413 A.V. Williams Bldg. #115, College Park, MD 20742-3275. E-mail: hjs@cs.umd.edu.

Manuscript received 28 June 2006; revised 2 Jan. 2007; accepted 23 Apr. 2007; published online 22 May 2007.

Recommended for acceptance by D. Fleet and B. Kimia.

For information on obtaining reprints of this article, please send e-mail to: tpami@computer.org, and reference IEEECS Log Number TPAMI-0476-0606. Digital Object Identifier no. 10.1109/TPAMI.2007.1182.

Authorized licensed use limited to: University of Durham. Downloaded on August 05, 2025 at 14:43:52 UTC from IEEE Xplore. Restrictions apply.
0162-8828/08/\$25.00 © 2008 IEEE Published by the IEEE Computer Society

“best-first”) rather than in a predetermined order, as in the depth-first method. In other words, at each step of the algorithm, the next nonobject element to be visited is the closest one to q , which has yet to be visited. This is achieved by storing the nonobject elements of the search hierarchy in a priority queue *Queue* according to this order. *Queue* is initialized to contain the root of the search hierarchy at a distance of 0 from q , and as nonobject elements are dequeued, their immediate descendants e that are nonobject elements are enqueued with their corresponding distances from q if $d(q, e) < D_k$, whereas immediate descendants o that are objects are inserted into L if $d(q, o) < D_k$, where D_k , the distance of the current k th-nearest neighbor of q , is initialized to ∞ . The algorithm repeatedly removes nonobject elements from *Queue* until it is empty or until encountering a nonobject element that is farther from q than D_k , at which time, it halts as it has found the k nearest neighbors, which are now in L . In order for the algorithm to be correct, the distance $d(q, e)$ of any nonobject element e from the query object q must be less than or equal to the distance from q to any object in e 's descendants [19]. Again, as in the depth-first method, this property is satisfied by letting $d(q, e)$ be MINDIST.

The drawback of the best-first method is that the priority queue may be rather large. In particular, the necessary amount of storage may be as large as the total number of nonobjects (and hence on the order of the number of objects) if the distance of each of the nonobjects from the query object q is approximately the same. In low dimensions, such an event is relatively rare as its occurrence requires two seemingly independent events—that is, that all objects lie in an approximate hypersphere centered at some point p and that the query object q be coincident with p . However, in high dimensions, where most of the data lies on the surface (for example, [5]) and the curse of dimensionality [3], [6] comes into play, and in metric spaces with concentrated distance histograms, this situation is less rare. In contrast, the amount of storage required by the depth-first method is bounded. In particular, it is proportional to the sum of k and the maximum depth of the search hierarchy, where in the worst case, all of the sibling nonobject elements must be retained for each partially explored nonobject element in the search hierarchy while executing the depth-first search.

Nevertheless, the advantage of the best-first algorithm over the depth-first algorithm is that it has been shown to be I/O optimal for $k = 1$ [4]. This means that the algorithm does not visit more than the minimum number of nonobject elements—that is, it avoids visiting nonobject elements that will eventually be determined to be too far from q due to poor initial estimates of D_k . This is equivalent to stipulating that the algorithm is range optimal, which means that the cost of finding the k nearest neighbors is the same as that of a range search with the search radius set to the distance from q to its k th-nearest neighbor.

As we point out above, the implementations of both the depth-first and best-first algorithms make heavy use of a lower bound MINDIST corresponding to the minimum distance at which a nearest object can be found via a vis the distance D_k to the current k th-nearest object from q . In this paper, we show how to also use an upper bound MAXNEARESTDIST corresponding to the maximum possible distance at which a nearest neighbor is guaranteed to be found in both the depth-first and best-first algorithms for finding the k nearest neighbors for arbitrary values of k . This upper bound was first

introduced in [22] for the case of $k = 1$ for the purpose of improving the initial estimate of D_k in the depth-first method. It was also proposed in [25] as an alternative to the MINDIST ordering for the processing of the nonobject immediate descendants of a nonobject element in the depth-first method but again limited to $k = 1$. However, for the purpose of ordering, MINDIST has been shown to be more useful [18], [25] and, hence, MAXNEARESTDIST is not discussed further here in this context. Therefore, in this paper we focus on the first purpose and show how to use the MAXNEARESTDIST upper bound for arbitrary values of k (that is, $k \geq 1$) to overcome the shortcomings that we pointed out earlier of both the depth-first (that is, pruning) and best-first (that is, size of the priority queue) k -nearest neighbor algorithms. Note that other upper bounds can be used in the k -nearest neighbor algorithms to yield what are termed probabilistically approximate nearest neighbors (for example, [9] and [11]), although they are beyond the scope of this paper.

The rest of this paper is organized as follows: Section 2 defines the MAXNEARESTDIST upper bound, whereas Section 3 describes how to use it in the process of finding the k nearest neighbors by incorporating it into the set L of the k nearest neighbors found so far, whereas Section 4 discusses the management of L in greater detail. Section 5 demonstrates how to incorporate it in a depth-first k -nearest neighbor algorithm to eliminate some elements from further consideration, whereas Section 6 demonstrates how to incorporate it in a best-first k -nearest neighbor algorithm to reduce the size of the priority queue of nonobject elements remaining to be processed. Section 7 presents the results of an experiment where the use of the MAXNEARESTDIST does indeed improve the performance of both the depth-first and best-first k -nearest neighbor algorithms, whereas concluding remarks are made in Section 8.

2 THE MAXNEARESTDIST UPPER BOUND

The key motivation for the introduction of the MAXNEARESTDIST upper bound in k -nearest neighbor finding have been the observations that until the first set of k candidate nearest neighbors has been found (which enables setting D_k to a value other than the initial value of ∞):

1. Regardless of the algorithm that is used, no visits of nonobject elements of the search hierarchy can be prevented.
2. In the best-first method, no insertion of nonobject elements into the priority queue can be avoided.

In fact, Larsen and Kanal [22] first introduced MAXNEARESTDIST as an alternative to the MAXDIST upper bound that was proposed by Fukunaga and Narendra [14] as a means of obtaining an upper bound on D_1 when finding the nearest neighbor using the depth-first method (see Fig. 1a). In particular, Fukunaga and Narendra's proposal for using MAXDIST assumed a very simple search hierarchy where objects are clustered into groups where cluster c has a cluster center M that need not necessarily correspond to an object in the cluster, and all objects in c lie within a distance r_{max} of M . Thus, in this context, the minimum distance at which the nearest neighbor could be found was indeed $\text{MAXDIST} = d(q, M) + r_{max}$.

Larsen and Kanal [22] improved upon Fukunaga and Narendra [14] by noting that once the cluster center is not

4 MANAGEMENT OF L

In this section, we discuss the management of the set L of the k nearest neighbors encountered so far using the principles outlined in Section 3. This discussion is independent of whether a depth-first or best-first algorithm is used to find the k nearest neighbors of query object q . In particular, we only assume that each time we process a nonobject element e of the search hierarchy, we insert in L all of e 's nonobject child elements along with their corresponding MAXNEARESTDIST values. In addition, before we attempt to insert the nonobject child elements of e into L , we remove e from L .

Given these assumptions about the new algorithms, the rest of this section is organized as follows: Section 4.1 presents the implementation of L . Section 4.2 proves a couple of key properties of L when used to store nonobject elements along with their corresponding MAXNEARESTDIST values in k -nearest neighbor finding. They are independent of which of the two variants of the original k -nearest neighbor algorithm is used. In particular, the only property of the algorithms that is used in the proofs is that they process the nonobject elements of the search hierarchy in increasing order of their MINDIST values either locally (depth-first algorithm) or globally (best-first algorithm). Section 4.3 describes in greater detail the process of inserting nonobject elements in L , whereas Section 4.4 discusses the process of the explicit removal of nonobject elements from L . The actual code for the two variants of the k -nearest neighbor finding algorithms that incorporate the MAXNEARESTDIST upper bound is given in Sections 5 and 6.

4.1 Implementation of L

The variants of the k -nearest neighbor algorithms that we describe need to be able to insert and remove specific objects and nonobjects with corresponding MAXNEARESTDIST values into and from the set L . In addition, we want to be able to access, as well as delete the farthest of the k nearest neighbors. We implement L using a priority queue as it enables the latter two operations to be performed without needless exchange operations as would be the case if L were to be implemented using an array. Each element e in L has two data fields E and D corresponding to the item i (object or nonobject) that e contains and i 's distance from q (that is, $d(q, i)$ or MAXNEARESTDIST(q, i), respectively), and a number of fields corresponding to control information specific to the data structure used to implement the priority queue (for example, a binary heap). We use the function MAXL(L) to access the element in L with the highest priority (that is, at the top of the queue or, equivalently, the first and most accessible element). When the queue L is full, then MAXL(L) corresponds to the k th-nearest neighbor (that is, the farthest of the known k nearest neighbors of q).

4.2 PROPERTIES OF L WHEN CONTAINING MAXNEARESTDIST VALUES

In this section, we prove some important properties of the modification of L described in Section 3. We assume that each object appears just once in the search hierarchy.² Associated with each nonobject element e in L is its

corresponding MAXNEARESTDIST(q, e) value. This value results from the postulation of the existence of an object o at this distance, and o is said to be *associated with e* , whether or not such an object actually exists. The key idea is that the particular positioning of o is what determines the MAXNEARESTDIST(q, e) value. In particular, we prove that each of the elements of L is associated with a unique object (Theorem 4.1) and that there is no need to insert a nonobject element e such that MAXNEARESTDIST(q, e) $\geq D_k$, which means that L contains at most k elements (Theorem 4.2).

Theorem 4.1. *The object o associated with each element e of L is unique.*

Proof. Since the objects in the set from which the neighbors are drawn are unique, once object o appears as one of the elements in L , it cannot be a member of one of the nonobject elements in L . Moreover, the fact that nonobject element e is removed from L before inserting any of e 's children into L ensures that no ancestor-descendant relationship exists between any two elements of L . Therefore, the object o associated with the nonobject element u of L at a distance of MAXNEARESTDIST is guaranteed to be unique even though it may not have been identified yet. In other words, at the time at which we are ready to insert u into L , its associated object o is not already in L nor is o associated with any other nonobject element in L .

Note that the definition of the MAXNEARESTDIST upper bound ensures that for each of the entries u in L , there is at least one object o in the data set whose maximum possible distance from q is the one that is associated with u . $\square$

Theorem 4.2. *There is no need to insert into L any nonobject element e such that MAXNEARESTDIST(q, e) $\geq D_k$, and thus the maximum size of L is k .*

Proof. There is no need to insert in L any nonobject element e such that MAXNEARESTDIST(q, e) $\geq D_k$ as e by itself cannot be used to further reduce the value of D_k . Moreover, failing to insert such an e in L does not affect the k -nearest neighbor finding process as regardless of the nature of the k -nearest neighbor finding algorithm that is used (that is, depth-first or best-first), the appropriate nonobject elements that are subsequently processed are explored in increasing order of their MINDIST values. Therefore, the fact that MINDIST(q, e) $\leq$ MAXNEARESTDIST(q, e) for all q and nonobject elements e means that if MINDIST(q, e) $\geq D_k$, then e will never be explored further anyway, and if MINDIST(q, e) $< D_k$, then e will be explored regardless of the value of MAXNEARESTDIST(q, e) and whether or not e was inserted into L . Therefore, not having inserted e in L makes no difference since the descendants of e will be explored anyway. The size of L is upper-bounded by k (actually, it can decrease) as when nonobject element e is removed from L , it could be the case that MAXNEARESTDIST(q, e_i) $\geq D_k$ for each of the immediate nonobject child elements e_i of e in which case none of them are inserted into L . Fig. 3 is an example that illustrates how the size of L can decrease for an element e with three spherical shell-like children e_a , e_b , and e_c , assuming a euclidean distance metric. In particular, in this case, we assume, without loss of generality, that D_k is equal to MAXNEARESTDIST(q, e) and that MAXNEARESTDIST(q, e_i) = $d(q, M_i) + r_{i, \min}$ for each of the

2. Search hierarchies where objects appear more than once as is the case for those based on a disjoint decomposition of the space from which the objects are drawn such as an R^+ -tree [28] and a PM quadtree [27] are more complex and beyond the scope of this paper.




Fig. 3. Example illustrating how the number of elements in L can decrease as a result of the situation that arises when the MaxNearestDist values of the three nonobject child elements e_a , e_b , and e_c of e_p are greater than the current value of D_k which is the MaxNearestDist value of e_p , assuming a euclidean distance metric.

children e_i ($i = \{a, b, c\}$). It is easy to see that $\text{MAXNEARESTDIST}(q, e_i) > D_k$ for each of e_i ($i = \{a, b, c\}$). $\square$

4.3 Insertion of Nonobject Elements into L

One of the key properties underlying any k -nearest neighbor algorithm is that the value of D_k , the distance of the k th-nearest neighbor of q , is nonincreasing.

This is ensured by Theorem 4.2 and the actual k -nearest neighbor algorithms (see Sections 5 and 6), which make use of procedure $\text{MAXNEARESTDISTINSERTL}$, given below, to update L as closer objects and nonobjects are found, thereby causing existing elements in L to be removed when L already contains k elements. Elements of L that are removed in such a manner are said to be removed *implicitly*, in contrast to elements that are removed *explicitly* whenever attempts are made to insert their children into L (see Section 4.4).

```

1 procedure MAXNEARESTDISTINSERTL( $e, s$ )
2 /* Insert element (object or nonobject)  $e$  at distance  $s$ 
   from query object  $q$  into the priority queue  $L$  using
   ENQUEUE. Assume that objects have precedence over
   nonobjects when they are at the same distance  $D_k$ 
   (that is, the  $k$ th-nearest neighbor) from the query
   object  $q$ .*/
3 if SIZE( $L$ ) =  $k$  then
4    $h \leftarrow \text{DEQUEUE}(L)$ 
5   if not ISOBJECT( $E(h)$ ) then
6     while not ISEMPY( $L$ )
7       and not ISOBJECT( $E(\text{MAXL}(L))$ )
8       and  $D(\text{MAXL}(L)) = D(h)$  do
9          $\text{DEQUEUE}(L)$ 
10    enddo
11  endif
12 endif
13 ENQUEUE( $e, s, L$ )
14 if SIZE( $L$ ) =  $k$  then
15   if  $D(\text{MAXL}(L)) < D_k$  then
16      $D_k \leftarrow D(\text{MAXL}(L))$ 
17   endif
18 endif
```

Procedure $\text{MAXNEARESTDISTINSERTL}$ proceeds as follows: If there are k candidate nearest neighbors in L (determined with the aid of the function $\text{SIZE}(L)$ that is not given here), then $\text{MAXNEARESTDISTINSERTL}$ precedes the

insertion, which is performed by ENQUEUE (not given here), by first dequeuing (that is, implicitly removing) the current farthest member (that is, the k th-nearest member) from L using DEQUEUE (not given here).³ Next, if there are k candidate nearest neighbors after the insertion, then $\text{MAXNEARESTDISTINSERTL}$ resets D_k to the distance of the current farthest nearest neighbor (see lines 15 and 16), accessed by the function MAXL (not given here, although in the case of a priority queue implementation of L as in this case, it is the current first element in the queue).

The situation is different when there are fewer than k candidate nearest neighbors in L as now there is no need to dequeue (that is, implicitly remove) any elements from L . An example of such a situation was given in the proof of Theorem 4.2 and is also always the case in the initial stages of the k -nearest neighbor finding process. Moreover, in this case, the insertion of an object or nonobject e into L does not cause a change in D_k as D_k indicates the minimum of all of the distance values that have been associated with the entry in L that corresponds to the k th-nearest neighbor of q , and D_k has not changed as a result of the current insertion of e at a distance $d(q, e) < D_k$ into L .

When L contains k elements and we must implicitly remove an element e in order to accommodate the insertion of the new element with distance less than D_k , then we must exercise some caution if there are several elements (objects and nonobjects) at the same distance D_k from q . The motivation for this is to try to minimize the need to search for e , and possibly not find e due to e having been implicitly removed from L earlier, when we subsequently attempt to explicitly remove e from L . This needless search is avoided by adopting some convention as to which element of L at distance D_k should be removed implicitly by $\text{MAXNEARESTDISTINSERTL}$ when there are several nonobjects in L having D_k as their MAXNEARESTDIST value, as well as objects at distance D_k .

We adopt the convention that objects have priority over nonobjects in the sense that in terms of nearness, objects have precedence over nonobjects in L . This means that when nonobjects and objects are at the same distance from q , the nonobjects appear closer to the maximum entry in the priority queue L (that is, $\text{MAXL}(L)$), which corresponds to the k th-nearest candidate neighbor. In particular, we stipulate that whenever insertion into a full priority queue results in dequeuing a nonobject element b with MAXNEARESTDIST value d , we check if the new $\text{MAXL}(L)$ entry c corresponds to a nonobject with the same MAXNEARESTDIST value d ; in which case, c is also dequeued. This loop continues until the new $\text{MAXL}(L)$ entry corresponds to an object at any distance including d or corresponds to a nonobject at any other distance $d' < d$, or L is empty (see lines 5-11). Note that D_k is only reset if exactly one entry has been dequeued (from a full priority queue) and the distance of the new $\text{MAXL}(L)$ entry is less than D_k (see lines 13-17). Otherwise, if we dequeue more than one entry, then even though the distance of the new $\text{MAXL}(L)$ entry may now be less than D_k , it cannot be used to reset D_k as L now contains fewer than k entries. In fact, it should be clear that D_k should not be reset as D_k has not been decreased since the only reason for the removal of the

3. Note the asymmetry between DEQUEUE that removes the item at the front of the queue while ENQUEUE inserts an item in its appropriate position in the queue. We also make use of a procedure REMOVEQUEUE in Section 4.4, which is the complement of ENQUEUE in that it removes a specific element from the queue, which may involve a search.

multiple nonobject entries is to avoid subsequent possibly needless searches when explicitly removing nonobject elements with MAXNEARESTDIST value D_k .

4.4 Removal of Nonobject Elements from L

As we pointed out in the proof of Theorem 4.2, there is no need for L to ever contain more than k elements. This simplifies the k -nearest neighbor algorithms considerably. However, it does mean that when we need to explicitly remove a nonobject element e from L just before inserting in L all of e 's nonobject child elements along with their corresponding MAXNEARESTDIST values that are less than D_k , it could be the case that e is no longer in L . This is because e may have been implicitly removed as a byproduct of the insertion of closer objects or nonobject elements as a result of their corresponding MAXNEARESTDIST values being smaller than that of e and thereby resulted in resetting D_k . Procedure MAXNEARESTDISTREMOVEL, given below, accomplishes this task, while also following our convention, set forth in Section 4.3, that objects have priority over nonobjects when they are in L at the same distance D_k (that is, the k th-nearest neighbor) from the query object q .

```

1 procedure MAXNEARESTDISTREMOVEL( $e$ )
2 /* Remove element (object or nonobject)  $e$  from the
   priority queue  $L$  using REMOVEQUEUE. Assume that
   objects have precedence over nonobjects when they are
   in  $L$  at the same distance  $D_k$  (that is, the  $k$ th-nearest
   neighbor) from the query object  $q$ . */
3 if MAXNEARESTDIST( $q, e$ ) <  $D_k$  or
4   (MAXNEARESTDIST( $q, e$ ) =  $D_k$  and
5      $D(\text{MAXL}(L)) = D_k$  and
6     not ISOBJECT( $E(\text{MAXL}(L))$ )) then
7   REMOVEQUEUE( $e, L$ )
8 endif
```

Procedure MAXNEARESTDISTREMOVEL proceeds as follows: When a nonobject e is to be removed explicitly from L and e 's MAXNEARESTDIST value is < D_k (see line 3), then e has to be in L as it is impossible for e to have been removed implicitly since D_k is nonincreasing in our algorithms (that is, the depth-first and best-first given in Sections 5 and 6). Therefore, we remove e and decrement the size of L using procedure REMOVEQUEUE, which is not given here. On the other hand, the situation is more complicated when e 's MAXNEARESTDIST value is equal to D_k (see line 4). First, if the maximum value associated with an element in L (that is, the one associated with $\text{MAXL}(L)$) is less than D_k (see line 5), then e cannot be in L , and we do not attempt to remove e . Such a situation arises, for example, when we have dequeued more than one nonobject due to having several nonobjects at distance D_k . Second, if the maximum value associated with an element in L (that is, the one associated with $\text{MAXL}(L)$) is equal to D_k , then there are two cases depending on whether the entry c in $\text{MAXL}(L)$ corresponds to an object or a nonobject (see line 6). If c corresponds to an object, then nonobject e cannot be in L as we have given precedence to objects, and all nonobjects at the same distance are either in L or they are all not in L . If c corresponds to a nonobject, then nonobject e has to be in L as all of the nonobjects at the same distance have been either removed implicitly together or retained, and in this case, by virtue of the presence of c in L , we know that they have been retained in L . Note that when we explicitly remove a nonobject at distance D_k from L , we do not remove all remaining nonobjects at the same distance from L

as this needlessly complicates the algorithm with no additional benefit as they will all be removed implicitly together later if at least one of them must be implicitly removed due to a subsequent insertion into a full priority queue.

5 DEPTH-FIRST ALGORITHM USING MAXNEARESTDIST

Incorporating the MAXNEARESTDIST upper bound in the depth-first k -nearest neighbor algorithm, thereby yielding what we characterize as a *maxnearest depth-first k -nearest neighbor algorithm*, is straightforward and is realized below by the recursive procedure MAXNEARESTDISTDF which is invoked with parameter e initialized to the root of the search hierarchy. In MAXNEARESTDISTDF, if the nonobject element e being visited is at the deepest level of the search hierarchy (usually referred to as a *leaf* or *leaf element*), then every object o in e that is nearer to q than the current k th-nearest neighbor of q (that is, $d(q, o) < D_k$) is inserted into L , with its associated distance from q (that is, $d(q, o)$), using procedure MAXNEARESTDISTINSERTL given in Section 2 (lines 2-9 of MAXNEARESTDISTDF). Otherwise (that is, e is not a leaf element), MAXNEARESTDISTDF generates the immediate successors of e , places them in a list $A(e)$, known as the *active list* of e , and proceeds to insert any of them whose MINDIST and MAXNEARESTDIST values are less than D_k into L (lines 13-21). It then proceeds to recursively process all of the nonobject child elements of e whose MINDIST value is less than D_k (line 24) after removing them from L if possible (line 26).

```

1 recursive procedure MAXNEARESTDISTDF( $e$ )
2 if ISLEAF( $e$ ) then /*  $e$  is a leaf with objects */
3   foreach object child element  $o$  of  $e$  do
4     Compute  $d(q, o)$ 
5     if  $d(q, o) < D_k$  or
6       ( $d(q, o) = D_k$  and SIZE( $L$ ) <  $k$ ) then
7       MAXNEARESTDISTINSERTL( $o, d(q, o)$ )
8     endif
9   enddo
10 else /*  $e$  is a nonleaf with nonobjects  $e_p$  */
11   Generate active list  $A$  with child elements  $e_p$  of  $e$ 
12   /*  $A$  is sorted in increasing order with respect to  $q$ 
   using MINDIST and processed in this order */
13   foreach element  $e_p$  of  $A$  do
14     /* Try to apply MAXNEARESTDIST while
       processing  $A$  in increasing order */
15     if MINDIST( $q, e_p$ ) >  $D_k$  then
16       exit_for_loop /* No further insertions */
17     elseif MAXNEARESTDIST( $q, e_p$ ) <  $D_k$  then
18       MAXNEARESTDISTINSERTL
19         ( $e_p, \text{MAXNEARESTDIST}(q, e_p)$ )
20     endif
21   enddo
22   foreach element  $e_p$  of  $A$  do
23     /* Process  $A$  in increasing order */
24     if MINDIST( $q, e_p$ ) >  $D_k$  then exit_for_loop
25   else
26     MAXNEARESTDISTREMOVEL( $e_p$ )
27     MAXNEARESTDISTDF( $e_p$ )
28   endif
29 enddo
30 endif
```

We now prove a couple of important properties of our algorithm. First, there are no nonobject elements in L when the algorithm (that is, MAXNEARESTDISTDF) terminates.

Theorem 5.1. *There are no nonobject elements in L when the maxnearest depth-first k -nearest neighbor algorithm (that is, MAXNEARESTDISTDF) terminates.*

Proof. We show this by contradiction. Suppose that L contains a nonobject element e upon termination. The fact that the algorithm has terminated means that $\text{MINDIST}(q, u) > D_k$ for all unprocessed nonobject elements u . However, the presence of e in L means that $\text{MAXNEARESTDIST}(q, e) \leq D_k$ and therefore by virtue of $\text{MINDIST}(q, e) \leq \text{MAXNEARESTDIST}(q, e)$, we know that e has been processed already, which means that e must have been removed explicitly from L , which contradicts our initial assumption that L contains nonobject elements upon termination. $\square$

Second, we prove our main result, which is that the number of nonobject elements that must be examined due to using the MAXNEARESTDIST upper bound is not increased.

Theorem 5.2. *The maxnearest depth-first k -nearest neighbor algorithm (that is, MAXNEARESTDISTDF) visits at most the same number of nonobject elements of the search hierarchy as the conventional depth-first algorithm, and may visit less.*

Proof. We know that D_k is nonincreasing as it is only updated when an object or nonobject at a distance less than D_k is inserted into L . This is true for both the conventional and maxnearest versions of the depth-first k -nearest neighbor algorithm. Inserting a nonobject element e into L in the maxnearest algorithm causes D_k to decrease or at the worst to maintain the same value. Suppose that D_k has indeed decreased so that it now has the value d_e instead of the previous value of d_p . This means that a nonobject element n with minimum distance d_n such that $d_e < d_n < d_p$ will not be visited, whereas n would have been visited had we not used the MAXNEARESTDIST upper bound, and thus the number of nonobject elements that are visited has decreased. $\square$

6 BEST-FIRST ALGORITHM USING MAXNEARESTDIST

Incorporating the MAXNEARESTDIST upper bound in the best-first algorithm, thereby yielding what we characterize as a *maxnearest best-first k -nearest neighbor algorithm*, is straightforward and is realized below by the procedure MAXNEARESTDISTBF, which is invoked with parameter e initialized to the root of the search hierarchy. Recall that in the depth-first algorithm, incorporation of MAXNEARESTDIST enabled the use of the nonobject elements of the search hierarchy to speed up the convergence of D_k to its final value thereby helping to prune the set of k candidate nearest neighbors instead of pruning only with the aid of the k nearest objects as in the standard implementation. It should be clear that both the fact that a best-first algorithm examines the nonobject elements in increasing MINDIST order and the fact that every nonobject element with MINDIST less than the final value of D_k must be examined together mean that no matter how fast the value of D_k converges to its final value, a best-first algorithm will never examine any extra nonobject elements. However, the use of MAXNEARESTDIST in the

best-first algorithm still helps to speed up the convergence of D_k to its final value which means that its use results in reducing the size of the priority queue *Queue* as fewer nonobject elements are inserted into it initially while D_k is at its initial value of ∞ .

```

1 procedure MAXNEARESTDISTBF( $e$ )
2   ENQUEUE( $e$ , MAXNEARESTDIST( $q$ ,  $e$ ),  $L$ )
3   ENQUEUE( $e$ , 0, Queue)
4   while not ISEMPY(Queue) do
5      $t \leftarrow$  DEQUEUE(Queue)
6      $e \leftarrow E(t)$ 
7     if  $D(t) > D_k$  then
8       return /* Found  $k$  nearest neighbors and exit */
9     else MAXNEARESTDISTREMOVEL( $e$ )
10    endif
11    if ISLEAF( $e$ ) then /*  $e$  is a leaf with objects */
12      foreach object child element  $o$  of  $e$  do
13        Compute  $d(q, o)$ 
14        if  $d(q, o) < D_k$  or
15          ( $d(q, o) = D_k$  and SIZE( $L$ )  $< k$ ) then
16          MAXNEARESTDISTINSERTL( $o$ ,  $d(q, o)$ )
17        endif
18      enddo
19    else /*  $e$  is a nonleaf */
20      foreach child element  $e_p$  of  $e$  do
21        if  $\text{MINDIST}(q, e_p) < D_k$  then
22          if  $\text{MAXNEARESTDIST}(q, e_p) < D_k$  then
23            MAXNEARESTDISTINSERTL
24              ( $e_p$ , MAXNEARESTDIST( $q$ ,  $e_p$ ))
25          endif
26          ENQUEUE( $e_p$ ,  $\text{MINDIST}(q, e_p)$ , Queue)
27        endif
28      enddo
29    endif
30  return

```

MAXNEARESTDISTBF processes all nonobject elements in *Queue* in the order in which they appear in *Queue* (that is, the element e at the front is processed first). We first remove e from *Queue* (lines 5-6) and also check if e should be explicitly removed from L using the same method as in MAXNEARESTDISTDF (line 9). Recall that this step ensures that the objects that are associated with the different entries in L are unique. This removal step is missing in a variant of a best-first k -nearest neighbor algorithm that uses MAXNEARESTDIST proposed by Ciaccia et al. for the M-tree [12] thereby possibly leading to erroneous results. Next, we check if e is a leaf element; in which case, we examine its constituent objects using the same method as in MAXNEARESTDISTDF (lines 12-18). Otherwise, we examine each of e 's child elements e_p and insert e_p and its associated $\text{MINDIST}(q, e_p)$ value into *Queue* (line 26) if $\text{MINDIST}(q, e_p)$ is less than D_k (line 21). When $\text{MINDIST}(q, e_p) \leq D_k$, we also check if e_p 's associated MAXNEARESTDIST(q, e_p) value is less than D_k (line 22); in which case, we use MAXNEARESTDISTINSERTL to insert e_p and MAXNEARESTDIST(q, e_p) into L and possibly reset D_k (line 23). As in MAXNEARESTDISTDF, this action may cause some elements (both objects and nonobjects) to be implicitly removed from L . Thus, the

MAXNEARESTDIST upper bound is used here to tighten the convergence of D_k to its final value.

Notice that in contrast to the depth-first algorithm (MAXNEARESTDISTDF), the nonobject child elements e_p of nonobject element e (that is, the elements of the active list of e) are not sorted with respect to their distance (MINDIST or MAXNEARESTDIST) from q before testing for the possibility of insertion into L and enqueueing into *Queue* (lines 20-28). In particular, assuming data of a fixed dimension and that the active list contains T elements, there is no advantage in incurring the extra time needed to sort the child elements (that is, $O(T \log T)$ time) since all that the sort can accomplish is avoiding the tests (that is, $O(T)$ time). In other words, unlike the depth-first algorithm, in the best-first algorithm, there is no need to worry about ordering the processing of the elements of the active list.

We now prove a couple of important properties of our algorithm. First, there are no nonobject elements in L when the algorithm (that is, MAXNEARESTDISTBF) terminates.

Theorem 6.1. *There are no nonobject elements in L when the maxnearest best-first k -nearest neighbor algorithm (that is, MAXNEARESTDISTBF) terminates.*

Proof. The fact that each time that a nonobject element e is removed from the priority queue *Queue*, e is also explicitly removed from L if it is in L by virtue of its MAXNEARESTDIST value being less than D_k , and the fact that $\text{MINDIST}(q, e) \leq \text{MAXNEARESTDIST}(q, e)$ together ensure that there are no nonobject elements left in L when the best-first algorithm terminates (that is, when the distance value associated with the first element in the priority queue *Queue* is greater than D_k), and thus, the elements in L are the k nearest neighbors of q . $\square$

Theorem 6.1 is of interest as when the best-first algorithm terminates, it is quite likely that the priority queue *Queue* is not empty as we do not constantly check it for the presence of nonobject elements with associated distance values greater than D_k each time the value of D_k decreases.

Second, we prove our main result which is that the maximum number of nonobject elements that may be in the priority queue due to using the MAXNEARESTDIST upper bound is not increased.

Theorem 6.2. *When the maxnearest best-first k -nearest neighbor algorithm (that is, MAXNEARESTDISTBF) terminates, the maximum size attained by the priority queue *Queue* is at most as large as that of the conventional best-first algorithm, and may be less.*

Proof. We know that D_k is nonincreasing as it is only updated when an object (nonobject) at a distance (MAXNEARESTDIST) less than D_k is inserted into L . This is true for both the conventional and maxnearest versions of the best-first k -nearest neighbor algorithm. Inserting a nonobject element e into L in the maxnearest algorithm causes D_k to decrease or at the worst to maintain the same value. Suppose that D_k has indeed decreased so that it now has the value d_e instead of the previous value of d_p . This means that a subsequently processed nonobject child element n with minimum distance d_n such that $d_e < d_n < d_p$ will not be inserted into *Queue*, whereas n would have been inserted into *Queue* had we not used the MAXNEARESTDIST upper bound, and thus the size of *Queue* has decreased. $\square$




Fig. 4. Block decomposition at the top 2 levels (depth 2 shown with darker borders than depth 1) in an R^* -tree for 100 points in a 512 x 512 space with the origin at the upper-left corner where the fanout at each node lies between 2 and 4. The maximum depth of the tree is 6. q1, q2, q3, q4, and q5 correspond to query points.

7 EXPERIMENTAL RESULTS

In this section, we demonstrate, with the aid of an example, situations where use of MAXNEARESTDIST can lead to additional pruning in the depth-first algorithm and likewise to a reduction in the size of the priority queue in the best-first algorithm. In particular, assuming a euclidean distance metric, we applied the depth-first and best-first algorithms with and without the use of MAXNEARESTDIST to the set of 100 two-dimensional points given in Fig. 4, stored in the R^* -tree [2], which is an object hierarchy where the minimum bounding boxes are hyperrectangles instead of spheres as is the case for the SS-tree [29]. The R^* -tree has the desirable property that overlap is kept low between minimum bounding boxes at the same level (that is, they are more likely to be disjoint or close to disjoint).

Fig. 5 shows the difference in performance of the depth-first (DF) and best-first (BF) algorithms with and without use of MAXNEARESTDIST for values of k , the number of nearest neighbors sought, ranging from 1 to 6 for five different query points labeled q1 ... q5 on the data in Fig. 4. From this example, we see that the improvements/savings are maximized as k gets smaller vis a vis the maximum capacity (that is, fanout) of the nodes of the search hierarchy that is used. This is not surprising as MAXNEARESTDIST, which is most effective at the initial stage of the search, cannot take effect until at least k nonobjects have been processed. Thus, when the capacity is less than k , it does not come into effect until nodes at depth 1 have been processed, which reduces its pruning power both in terms of nodes to be processed (depth-first) and enqueued (best-first).

We also observe that the benefit of using MAXNEARESTDIST is more pronounced in the case of the best-first algorithm where one of our examples demonstrated a 400 percent improvement while 10-15 percent seems a more reasonable expectation for the best-first algorithm. The fact that MAXNEARESTDIST is more effective in the best-first algorithm than the depth-first algorithm is attractive as the best-first algorithm is I/O optimal and thus we are overcoming its only drawback, which is the potentially large queue size.

In general, from our example, it can be seen that the improvement/savings that can be obtained from use of

Number of Neighbors	Query Points									
	q1=(335,453)		q2=(339,343)		q3=(170,309)		q4=(190,365)		q5=(326,353)	
	DF	BF	DF	BF	DF	BF	DF	BF	DF	BF
1	15:11	12:3	13:9	13:5	9:9	8:5	9:9	8:4	17:13	18:5
2	15:11	12:3	13:9	13:5	9:9	8:5	9:9	8:4	17:13	18:5
3	21:15	12:6	21:19	13:10	13:13	10:5	11:9	10:6	21:19	18:8
4	26:19	12:9	23:20	13:10	15:13	11:7	13:11	10:6	23:19	18:11
5	26:24	12:9	23:21	17:13	16:14	11:7	14:11	10:7	23:21	18:14
6	26:25	12:10	24:23	17:13	16:14	11:10	16:14	10:7	26:23	18:16

Fig. 5. The effect of using MaxNearestDist in the depth-first and best-first k -nearest neighbor finding algorithm for five different query points in the 100 data point R^* -tree in Fig. 4, where each node contains at least 2 and at most four points. Entries A:B in columns labeled DF indicate the number of recursive calls to the procedure using MaxNearestDist (B) and without using it (A). Entries C:D in columns labeled BF indicate the maximum size of the priority queue Queue using MaxNearestDist (D) and without using it (C).

MAXNEARESTDIST are heavily dependent on the underlying distribution of the data and on the positioning of the query point. However, most importantly, the performance of the algorithms can only be improved by using MAXNEARESTDIST, whereas this is not the case in some of the prior usage of MAXNEARESTDIST where it was used to order the processing of the elements of the active list in the depth-first algorithm (for example, [13], [25]).

8 CONCLUDING REMARKS

We have shown how to use MAXNEARESTDIST, an upper bound corresponding to the maximum possible distance at which a nearest neighbor is guaranteed to be found to enhance the performance of both a depth-first branch and bound and a best-first k -nearest neighbor finding algorithm by virtue of yielding tighter initial estimates of D_k , the distance at which the k th-nearest neighbor is found. This enables us to start pruning elements of the search hierarchy (both objects and nonobjects) in the depth-first algorithm and to avoid entering nonobject elements in the priority queue in the best-first algorithm. Thus, we see that use of MAXNEARESTDIST enhances the performance of both the depth-first and best-first algorithms by addressing their shortcomings—that is, reducing the number of nonobject elements that need to be examined by the former and reducing the storage requirements of the latter at no extra cost in the execution time of the latter and no extra storage requirements for the former. Nevertheless, it is important to bear in mind that we are not saying anything about the relative performance of the two algorithms, which is a more general issue and beyond the scope of this paper.

Some implementations of the best-first nearest neighbor algorithm (for example, [16], [17], [18], [24]) also store the objects in a priority queue, thereby enabling the algorithms that employ this method to be incremental. This means that now both the objects and nonobjects are visited in increasing order of their distance from q , and the objects are also reported in increasing order of their distance from q . They are designed for the case that k is not known in advance, thereby making them inappropriate for use with the MAXNEARESTDIST upper bound as no nonobject

elements can be excluded from *Queue* since they may all be eventually needed should k get sufficiently large (the same holds for probabilistic algorithms such as [9]).

Note, that the complexity of the process of computing the MAXNEARESTDIST upper bound depends on the nature of clustering process used to form the search hierarchy, as well as the domain of the data. For example, in d dimensions, using the euclidean distance metric, its complexity is the same as that of MINDIST when the cluster elements are minimum bounding hyperspheres (that is, $O(d)$), whereas when the cluster elements are minimum bounding hyperrectangles, the complexity of MAXNEARESTDIST is $O(d^2)$ while the complexity of MINDIST in this case is just $O(d)$.

The utility of the MAXNEARESTDIST upper bound depends on the distribution of the underlying data and also on the nature of the clustering methods that are applied in forming the search hierarchies. An interesting and open question is determining the type of data distributions and clustering methods for which MAXNEARESTDIST is most effective. For example, MAXNEARESTDIST may be most useful when using nonstandard clustering methods where objects are not necessarily associated with the closest cluster center (see object o in Fig. 6). Similarly, as another example, consider clusters formed by the five interlocking Olympic rings. On the other hand, MAXNEARESTDIST is not particularly useful for uniformly distributed data or when the query object is inside one of the clusters.




Fig. 6. Examples illustrating that an object o is not necessarily associated with the element whose cluster center is closest to o . In particular, in the case of each of parts (a), (b), and (c), o is associated with element e_a , whose cluster center M_a is farther away from o than cluster center M_b of element e_b .

ACKNOWLEDGMENTS

This work was supported in part by the US National Science Foundation under grants EIA-99-00268, EIA-99-01636, IIS-0086162, EIA-0091474, and CCF-0515241, Microsoft Research, and the University of Maryland General Research Board.

REFERENCES

- [1] S. Arya, D.M. Mount, N.S. Netanyahu, R. Silverman, and A.Y. Wu, "An Optimal Algorithm for Approximate Nearest Neighbor Searching in Fixed Dimensions," *J. ACM*, vol. 45, no. 6, pp. 891-923, Nov. 1998.
- [2] N. Beckmann, H.-P. Kriegel, R. Schneider, and B. Seeger, "The R*-Tree: An Efficient and Robust Access Method for Points and Rectangles," *Proc. ACM SIGMOD Conf.*, pp. 322-331, June 1990.
- [3] R.E. Bellman, *Adaptive Control Processes*. Princeton Univ. Press, 1961.
- [4] S. Berchtold, C. Böhm, D.A. Keim, and H.-P. Kriegel, "A Cost Model for Nearest Neighbor Search in High-Dimensional Data Space," *Proc. 16th ACM SIGACT-SIGMOD-SIGART Symp. Principles of Database Systems*, pp. 78-86, May 1997.
- [5] S. Berchtold, C. Böhm, and H.-P. Kriegel, "Improving the Query Performance of High-Dimensional Index Structures by Bulk-Load Operations," *Advances in Database Technology-Proc. First Int'l Conf. Extending Database Technology*, H.-J. Schek, F. Saltor, I. Ramos, and G. Alonso, eds., Springer-Verlag, Lecture Notes in Computer Science, vol. 1377, pp. 216-230, Mar. 1998.
- [6] K.S. Beyer, J. Goldstein, R. Ramakrishnan, and U. Shaft, "When Is 'Nearest Neighbor' Meaningful?" *Proc. Seventh Int'l Conf. Database Theory*, C. Beeri and P. Buneman, eds., Springer-Verlag, Lecture Notes in Computer Science, vol. 1540, pp. 217-235, Jan. 1999.
- [7] C. Böhm, S. Berchtold, and D.A. Keim, "Searching in High-Dimensional Spaces: Index Structures for Improving the Performance of Multimedia Databases," *ACM Computing Surveys*, vol. 33, no. 3, pp. 322-373, Sept. 2001.
- [8] A.J. Broder, "Strategies for Efficient Incremental Nearest Neighbor Search," *Pattern Recognition*, vol. 23, nos. 1-2, pp. 171-178, Jan. 1990.
- [9] B. Bustos and G. Navarro, "Probabilistic Proximity Searching Algorithms Based on Compact Partitions," *J. Discrete Algorithms*, vol. 2, no. 1, pp. 115-134, Mar. 2004.
- [10] E. Chávez, G. Navarro, R. Baeza-Yates, and J. Marroquín, "Searching in Metric Spaces," *ACM Computing Surveys*, vol. 33, no. 3, pp. 273-322, Sept. 2001.
- [11] P. Ciaccia and M. Patella, "PAC Nearest Neighbor Queries: Approximate and Controlled Search in High-Dimensional and Metric Spaces," *Proc. 16th IEEE Int'l Conf. Data Eng.*, pp. 244-255, Feb. 2000.
- [12] P. Ciaccia, M. Patella, and P. Zezula, "M-Tree: An Efficient Access Method for Similarity Search in Metric Spaces," *Proc. 23rd Int'l Conf. Very Large Data Bases*, M. Jarke, M.J. Carey, K.R. Dittrich, F. H. Lochovsky, P. Loucopoulos, and M.A. Jeusfeld, eds., pp. 426-435, Aug. 1997.
- [13] H. Ferhatosmanoglu, I. Stanoi, D. Agrawal, and A. El Abbadi, "Constrained Nearest Neighbor Queries," *Advances in Spatial and Temporal Databases-Proc. Seventh Int'l Symp.*, C.S. Jensen, M. Schneider, B. Seeger, and V.J. Tsotras, eds., Springer-Verlag, Lecture Notes in Computer Science, vol. 2121, pp. 257-278, July 2001.
- [14] K. Fukunaga and P.M. Narendra, "A Branch and Bound Algorithm for Computing k -Nearest Neighbors," *IEEE Trans. Computers*, vol. 24, no. 7, pp. 750-753, July 1975.
- [15] V. Gaede and O. Günther, "Multidimensional Access Methods," *ACM Computing Surveys*, vol. 20, no. 2, pp. 170-231, June 1998.
- [16] A. Henrich, "A Distance-Scan Algorithm for Spatial Access Structures," *Proc. Second ACM Workshop Geographic Information Systems*, N. Pissinou and K. Makki, eds., pp. 136-143, Dec. 1994.
- [17] G.R. Hjaltason and H. Samet, "Ranking in Spatial Databases," *Advances in Spatial Databases-Proc. Fourth Int'l Symp.*, M.J. Egenhofer and J.R. Herring, eds., Springer-Verlag, Lecture Notes in Computer Science, vol. 951, pp. 83-95, Aug. 1995.
- [18] G.R. Hjaltason and H. Samet, "Distance Browsing in Spatial Databases," *ACM Trans. Database Systems*, vol. 24, no. 2, pp. 265-318, June 1999.
- [19] G.R. Hjaltason and H. Samet, "Incremental Similarity Search in Multimedia Databases," Computer Science Technical Report TR-4199, Univ. of Maryland, Nov. 2000.
- [20] G.R. Hjaltason and H. Samet, "Index-Driven Similarity Search in Metric Spaces," *ACM Trans. Database Systems*, vol. 28, no. 4, pp. 517-580, Dec. 2003.
- [21] B. Kamgar-Parsi and L.N. Kanal, "An Improved Branch and Bound Algorithm for Computing k -Nearest Neighbors," *Pattern Recognition Letters*, vol. 3, no. 1, pp. 7-12, Jan. 1985.
- [22] S. Larsen and L.N. Kanal, "Analysis of k -Nearest Neighbor Branch and Bound Rules," *Pattern Recognition Letters*, vol. 4, no. 2, pp. 71-77, Apr. 1986.
- [23] C. Merkwirth, U. Parlitz, and W. Lauterborn, "Fast Exact and Approximate Nearest Neighbor Searching for Nonlinear Signal Processing," *Physical Rev. E (Statistical Physics, Plasmas, Fluids, and Related Interdisciplinary Topics)*, vol. 62, no. 2, pp. 2089-2097, Aug. 2000.
- [24] G. Navarro, "Searching in Metric Spaces by Spatial Approximation," *Vldb J.*, vol. 11, no. 1, pp. 28-46, Aug. 2002.
- [25] N. Roussopoulos, S. Kelley, and F. Vincent, "Nearest Neighbor Queries," *Proc. ACM Sigmod Conf.*, pp. 71-79, May 1995.
- [26] H. Samet, *Foundations of Multidimensional and Metric Data Structures*. Morgan-Kaufmann, 2006.
- [27] H. Samet and R.E. Webber, "Storing a Collection of Polygons Using Quadrees," *ACM Trans. Graphics*, vol. 4, no. 3, pp. 182-222, July 1985.
- [28] M. Stonebraker, T. Sellis, and E. Hanson, "An Analysis of Rule Indexing Implementations in Data Base Systems," *Proc. First Int'l Conf. Expert Database Systems*, pp. 353-364, Apr. 1986.
- [29] D.A. White and R. Jain, "Similarity Indexing with the SS-Tree," *Proc. 12th IEEE Int'l Conf. Data Eng.*, S.Y.W. Su, ed., pp. 516-523, Feb. 1996.
- [30] P. Zezula, G. Amato, V. Dohnal, and M. Batko, *Similarity Search: The Metric Space Approach*. Springer, 2006.



Hanan Samet received the BS degree in engineering from the University of California, Los Angeles, and the MS degree in operations research and the MS and PhD degrees in computer science from Stanford University, Stanford, California. In 1975, he joined the Computer Science Department, University of Maryland, College Park, where he is now a professor. He is a member of the Computer Vision Laboratory of the Center for Automation

Research and also has an appointment at the University of Maryland Institute for Advanced Computer Studies. At the Computer Vision Laboratory, he leads a number of research projects on the use of hierarchical data structures for geographic information systems. His research group has developed the QUILT system, which is a GIS based on hierarchical spatial data structures such as quadrees and octrees, the SAND system which integrates spatial and nonspatial data, the SAND Spatial Browser, which enables browsing through a spatial database using a graphical user interface, the VASCO spatial indexing applet (<http://www.cs.umd.edu/~hjs/quadtree/index.html>), and a symbolic image database system. His research interests are data structures, computer graphics, geographic information systems, computer vision, robotics, and database management systems. He is the author of the recent book titled *Foundations of Multidimensional and Metric Data Structures* (<http://www.mkp.com/multidimensional>), published by Morgan-Kaufmann, an imprint of Elsevier, in 2006, and of the first two books on spatial data structures titled *Design and Analysis of Spatial Data Structures*, and *Applications of Spatial Data Structures: Computer Graphics, Image Processing, and GIS*, both published by Addison-Wesley in 1990. He is an area editor of *Graphical Models and Image Processing* and on the editorial board of *Image Understanding*, *Journal of Visual Languages*, *Geoinformatica*, and *Journal of Spatial Cognition and Computation*. He is a fellow of the IEEE, the ACM, and International Association for Pattern Recognition (IAPR).

► For more information on this or any other computing topic, please visit our Digital Library at www.computer.org/publications/dlib.